

A BRAG Guide to HPC

Ethan Goan



Contents

Preface	3
1 Introduction and Preliminaries	5
1.1 Getting HPC Access	5
1.1.1 Mac and Linux	5
1.1.2 Windows	6
1.2 Once Logged In	6
1.3 Transferring Files to HPC	7
2 Submitting Jobs	8
2.1 Interactive Jobs	8
2.2 Modules	9
2.3 Installing R Packages	11
2.4 Submitting Batch Jobs	11
2.4.1 Installing More Packages	13
2.4.2 Quick Note on Python	14

3	Examples	14
3.1	Multicore Processing - Let the OS do the Hard Work	14
3.1.1	Tips for bulk Submitting Jobs	17
3.2	Multicore Processing - The Hard Way	18
4	Quick Guide	19
5	Other Tips	20

Preface

As the size of data sets increase and models become increasingly complex, the need for High Performance Computing (HPC) systems becomes increasingly prominent. The purpose of this document is to provide a gentle introduction to using the High Performance Computing (HPC) cluster from a statistical perspective. Sample scripts and examples are provided, and intended to serve as a basis for any future work people may need. From these examples provided in this guide, you will be able to modify them to your need to make submitting jobs to HPC as easy as possible. It is hoped that this document can serve as a reference guide, a first port of call to those who might not be familiar with HPC and want to benefit from the facilities. It is also important to note that the scheduler used at QUT is used extensively at many universities and research facilities, so knowledge of QUT's HPC system will extend much past QUT.

This guide will start with a brief introduction to UNIX like systems, and how to navigate a command line. The guide will also show how you can bypass much of the command line use, to remove as much prerequisite knowledge of UNIX as possible. The HPC scheduler is then introduced, along with modules currently installed on the cluster and how to access them. Sample scripts are provided with instructions on how to submit both batch and interactive jobs. Instructions on how to customise your personal environment and install your own packages is also provided. This guide will primarily focus on programs written in R, though instructions can be followed with only slight modifications for other platforms and languages such as MATLAB, Python and Mathematica.

Additional tips and tricks to help get the most out of the cluster are also provided. The tips include suggestions on how to (hopefully) simplify your HPC experience. This will be demonstrated through sample scripts that you can run yourself.

From my experience, others have often already encountered and solved many of the problems. I have also noticed that more often than not, others with more experience have developed far more elegant solutions to the same problems I have faced. It is for this reason that I have decided to make this document, to impart some of the knowledge

that others with considerably more experience have passed on to me. In this spirit, I am happy for this to be a living document that all can edit, to provide your own tips and tricks to solve potential problems you think others might encounter. Through collaboration, we can all benefit from each others work and streamline our development process, by not needlessly addressing a problem others have already solved.

1 Introduction and Preliminaries

The QUT HPC facilities provide access to,

- 212 compute nodes
- 3780 Intel Xeon Cores
- Approx. 200G B RAM per Compute Nodes
- 34 TB of main storage
- 1800 TB additional storage in file store
- 24 Tesla GPUs
- Visualisations services and more

1.1 Getting HPC Access

Before accessing HPC, you will first need to apply for access. This can be done via HiQ by following the instructions on this link [here](#). Access will usually be granted in a couple of days at the most.

Once access has been granted, you will need a program to access HPC. If you are using a Linux or Mac, you won't need to install anything, though if you are using Windows, it is recommended to use the program PuTTY [1]. Instructions on logging in to HPC for each OS are given here.

1.1.1 Mac and Linux

To log into HPC, we will be using the Secure SHell (SSH) protocol, which we can access through the command line on Mac and Linux devices. To login, first open a command line prompt (terminal) in your computer. To login, you will simply need to type in,

```
ssh <your_qut_id>@lyra.qut.edu.au
```

where you will replace <your_qut_id> with your QUT username, ie.

```
#if you applied for HPC access with staff account
ssh jane.doe@lyra.qut.edu.au

#if you applied for HPC access with student account
ssh n12345678@lyra.qut.edu.au
```

1.1.2 Windows

To login on windows, first install PuTTY by following the instructions [here](#). Once downloaded, simply open the putty window and enter `lyra.qut.edu.au` into the bar as shown in Figure 1. You will then have to enter your QUT username and password.

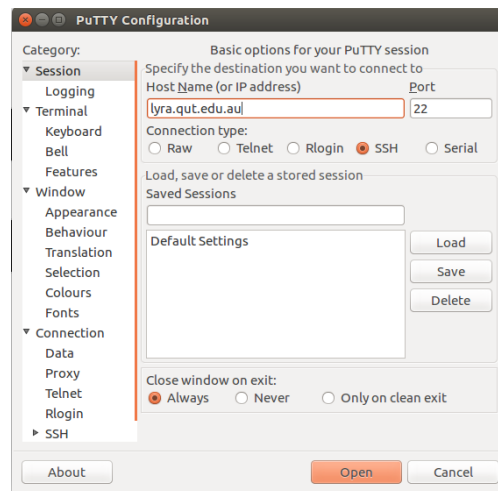


Figure 1: Example PuTTY window

1.2 Once Logged In

After following the previous commands you will be logged into the head node of the Lyra HPC cluster. Though you will now be logged in, you have not been allocated any computational resources yet. The head node is available to all users once logged

in, and is designed for performing only simple tasks, such as text editing or checking on currently running jobs.

1.3 Transferring Files to HPC

Now that you have access to HPC, you will want to transfer some code over so you can start running some jobs. There are a few ways to do this, though this guide will cover the easier/most prominent ways to do this.

To get your code onto HPC, I would recommend simply using Git and cloning your repo onto HPC. If you are unsure on how to use Git, I would strongly recommend you take some time to learn some basics and start using it for your software projects. Knowledge of Git is not required for this guide, but it will make your software development life significantly better, and allow you to distribute your research much easier.

You can clone your repo to HPC via the terminal that is connected to HPC by simply running the command,

```
#change to the home directory
git clone <remote_url_for_your_repo>
```

It is also helpful to be able to transfer other files to and from HPC. The easiest way to do this is to mount a network directory, so you can copy and paste files to and from HPC just like you would on your desktop/laptop. Instructions on how to do this for Windows, Mac and Linux using SSHFS is provided [here](#). After you install the required programs, you can mount the directory creating a folder on your desktop where you want to mount your HPC files by running the following command,

```
#all of this in a single line/command
sshfs -o allow_other <your_qut_username>@lyra.qut.edu.au
:/home/<your_qut_username> <
path_to_mount_location_on_desktop>
```

2 Submitting Jobs

Whilst it may seem like ample resources are available, they are finite and accessed by many people, thus access to these services needs to be managed. Access to computing resources is scheduled by the Portable Batch System (PBS). Before you run any software, you must tell the PBS manager what resources you require and how long you will need them for. This is done by submitting jobs.

There are two main types of jobs you can submit to HPC; interactive and batch jobs. Interactive jobs provide you with an active terminal similar to what you will currently use on your desktop. Interactive jobs are useful for debugging and ensuring all of your code can run.

Interactive jobs are often convenient, though you need to have your interactive session open for the job to continue. Batch jobs are intended for jobs that will need to run for longer, or when multiple jobs need to be submitted. Unlike interactive jobs, after you submit a batch job, you can close your connection to HPC altogether, go get a coffee, put your feet up and relax while your job processes on HPC in the background.

We will first start with interactive jobs, and then show how we can move our work to batch jobs later on.

2.1 Interactive Jobs

To submit a job to HPC, we use the `qsub` command, along with a few arguments to tell the scheduler what resources we require, and how long we need them for. To submit an interactive job, run the following command,


```
#submit an interactive job to HPC
#replace the HH:MM:SS with the amount of time you
#expect your job to run
#Replace XXX with the amount of RAM you need
#Replace YYY with the number of CPUS you need
#Replace ZZZ with the name of your job
#you can call it whatever you like :)
qsub -I -S /bin/bash -l walltime=HH:MM:SS,mem=XXXg,ncpus=
YYY -N ZZZ
```

Change the variables supplied here with the time and resources required for your work. After running this command, you will have to wait for your requested resources to be allocated to you. The time taken for your job to be accepted will depend on the amount of resources you requested. If you asked for 8 GB of RAM, 2 CPUs for only a couple hours, your job should be accepted within a minute or so. If you request 100 GB, 20 CPUs for 12 days, expect to wait a very long time for your interactive job to be accepted¹

2.2 Modules

Now that you have transferred your code across over to HPC and have been allocated resources for a job, you can start loading and installing the required packages you need. HPC has many programs already installed, though they aren't initially loaded when you log in. These pre-installed programs are stored as *modules* that need to be first loaded before you can use them. To see the modules currently installed on HPC, run the command,

```
#see what modules are available
module avail
```

¹You can request these resources, but will need to submit a batch job.

From the output of this, you may begin to appreciate why not all of the packages are loaded on startup, there is an awful lot of them. You can search through the output to find any modules you are interested in. Once you have found the module you are interested in, you can load it with the `module load` command. An example of common modules that might be helpful are listed here.

```
#load in R
module load atg/R/3.4.1-foss-2016a

#load MATLAB
#many different versions available
#only need to load one you need
module load matlab/2016a
module load matlab/2016b
module load matlab/2017b

#load mathematica
module load mathematica/11.2.0-linux-x86_64

#load Python
#again many different versions available
module load python/2.7.13-foss-2017a-foss
module load python/3.5.1-foss-2016a
python/3.6.4-intel-2017a
```

For this guide, we will use R as an example, though you can adapt it for other programming languages with only small modification. So first we load in the R module with `module load atg/R/3.4.1-foss-2016a`. Once loaded, you can start R by simply typing R into the terminal.

2.3 Installing R Packages

If you try and run an R script now, you will likely find that it will throw an error saying that a package isn't available. The module we loaded before was the base R module, and unfortunately there aren't many R packages pre-installed on the Lyra cluster. This isn't a major limitation, we just need to install them ourselves. An pre-made R script has already been made, `install_r_packages.R`, which will install many of the common packages you will need. After you have cloned the example repository listed earlier, you can run the script to install all of the base packages with these commands,

```
#change to home directory
cd ~
git clone https://github.com/ethangoan/hpc_guide
#change directory to the repo
cd hpc_guide
#now run the install script
Rscript install_r_packages.R
```

This will take a while to run (a few hours I think), so you can either leave your terminal open and let the program run, or you can use the instructions in the next section, where we will learn to submit a batch job that will install all of the packages for you.

2.4 Submitting Batch Jobs

In the previous example, we saw how to submit an interactive job, load in R modules and install some base packages in an interactive session. We can achieve this same result by submitting a batch job, which will run on HPC without us having to intervene and leave the terminal open. Batch jobs are useful for programs that require a long time to run, since we can simply submit them and then forget about them (while they running at least).

Like submitting an interactive job, we need to specify the time and computational resources we require. Unlike interactive jobs, we specify these requirements through

a configuration file. In the guide repository, an example batch configurations script called `batch_jobs/install_packages_batch.sh` is supplied. This is a Bash script that is interpreted by the PBS scheduler, and specifies our requirements and which program we want to run. Computational requirements are listed at the top of the file in the commented out section. These are called the PBS directives.

```
#!/usr/bin/env bash

#PBS -N install_packages
#PBS -l ncpus=1
#PBS -l mem=2GB
#PBS -l walltime=20:00:00
#PBS -o install_packages_stdout.out
#PBS -e install_packagesstderr.out
```

The main differences here is the first line which is called the Shebang. This MUST be there in any batch configuration script, you will never need to change it. The other differences is the last two lines, which specifies where standard output and error messages will be written to.

Further down in the script you will see helper functions that will load all of the modules we need (for this example we only need the R module) and a function which invokes the R script to install the packages we need. These helper functions are called at the end of the script when the job has been submitted and accepted by the scheduler. We can submit this job using the following command,

```
#clone the guide repo into your home directory
#if you haven't already
cd ~
git clone https://github.com/ethangoan/hpc_guide
#change into repo directory
cd ~/hpc_guide
#change to the directory where the config file is
cd ./batch_jobs
#submit the batch job with qsub
qsub install_packages_batch.sh
```

Once you have submitted the job, you can track all of your submitted jobs using the command,

```
watch -n 1 qstat <your_QUT_username>
```

This will give you information on all the jobs you have submitted. You will be able to see whether they have commenced running, or if they are still running and how long they have been running for. Once the program has finished running, you can view the output of the installation script with the `cat` command.

```
#check the output of the program
cat install_packages_stdout.out
#check the error log to see if anything went wrong
cat install_packages_stderr.out
```

2.4.1 Installing More Packages

While running this installation script will install many of the most common packages, it is unlikely that it will install everything you require. To install more packages, I would suggest modifying the `instal_r_packages.R` script to include packages to want to install. There is a slight difference to installing packages when compared with

a typical desktop machine you own. Since you won't have root/administrator access on HPC, you will need to install the packages locally. The `install_r_packages.R` installs the packages locally and sets the relevant path variables so that R can find the packages we installed. To install more packages, simply edit the `packages` vector in that script and resubmit the batch job using the same commands as before.

2.4.2 Quick Note on Python

While there aren't many pre-installed packages for R on HPC, there is many for Python. Popular packages such as Numpy, Matplotlib, Scipy, Sklearn etc. are already installed and have their own module. To find these modules, simply run the `module avail` command and search for the module you require. Then load the module using the same `module load` command used previously.

3 Examples

Now we are set up on HPC and have installed some of the packages we need, we will go through some examples on how to get the most out of HPC.

3.1 Multicore Processing - Let the OS do the Hard Work

Many times when we want to process a large data set, we want to do a single task to each element in the data set, and sometimes this individual operation can be computationally expensive. An example is preprocessing all images in a large data set to remove certain artefacts, convert to a more convenient format etc.. It would be beneficial to process many of these items in the many available CPU cores on HPC. One method is to write a multi-threaded/multi-process script (not a simple task in R) to process the data. Another and far easier way to handle this is to create a script that processes a single item in the data set, and submit this job many times to the HPC cluster with a different observation from the data set as an input example. The idea is to let the OS and the scheduler do the hard part of organising multicore processing.

Another scenario when this type of processing is helpful is for model validation. Consider a case where you are commencing work on a new project, with a new type of data set and you want to run some experiments to bench mark the performance of different models with different parameters. For example, say you are fitting a mixture model, and you want to investigate the performance of the model for different number of mixtures, or a boosted regression tree, where you want to see how the accuracy of your predictions change when altering the parameters of the model. This type of experimentation can greatly benefit from this type of parallel processing, where you want to run several independent experiments with varying parameters. An example of how to do this is supplied in the `bt_examples` directory of the repo for this guide, where we will fit many different boosted regression tree models to try and predict the presence of breast cancer based on biopsy information [2].

When looking at the contents of the `bt_examples` directory, you will find a single R script `breast_cancer_bt.R` and ten batch scripts `batch_bt_XXXX.sh`. Each of these batch scripts will invoke the `breast_cancer_bt.R` script, though each script will supply a different command line argument to specify the number of trees we want to use in the model. We can submit batch jobs for all of these scripts using the following commands,

```
#Change to the bt_examples directory
cd ~/hpc_guide/bt_examples/
#now lets send all the batch scripts to qsub
#so we can submit jobs for them
#
#We can use a for loop to submit all scripts
#that end in .sh
for sub in $( ls /*.sh ); do qsub $sub; done
```

After running this command, you will see that ten independent jobs have been submitted. You can track the progress of these jobs by again running the command,

```
watch -n 1 qstat -u <your_QUT_username>
```

Once all of the jobs are completed, you will see a number of log files have been created, a standard output and a standard error file for each job submitted. You can view the output of these files using the `cat` command.

```
#view the output of a single job
cat bt_10000_stdout.out
cat bt_10000_stderr.out
#if the output file is long, you can display it
#in a slightly nicer format where you can scroll through
#using enter or the space bar
cat bt_10000_stdout.out | more
```

After running these jobs, you may want to remove the current stash of output files. You can do this using the `rm` command, though this should be used with caution. This command will remove files for good, and unlike your desktop system, after you remove a file in a UNIX like system, it is gone for good! This is another reason why you should use version control systems such as Git with remote back ups. If you accidentally delete all of your source files and you haven't backed them up using Git or something similar, they will be gone forever and you will have to start again from scratch!.

I stress the importance/danger associated with using the `rm` command to hopefully help you avoid disaster. In saying that, if used properly it is a simple and extremely useful command.

If you have named output files in the standard I have used throughout my examples (output scripts ending with `stdout.out` and `stderr.out`), then you can delete all of these files with the following commands,

```
#delete any files that end with stdout.out
rm ./*stdout.out
#delete any files that end with stderr.out
rm ./*stderr.out
```


3.1.1 Tips for bulk Submitting Jobs

Bulk submitting jobs in this way relies on you designing your original code to handle command line arguments. Adopting this type of program design is extremely helpful during experimentation, and is a part of general good coding practices. Don't hard-code anything you think may even have the slightest possibility of ever changing. Command line arguments are a great way to develop software that is highly modular, and generally easy to use (as long as you document what you have done!).

For bulk submitting jobs in this way, I would recommend using the bash scripts I have provided as a template, and simply modifying them to suit your needs. The components that you will need to modify include the PBS directives at the top that define your computational requirements. Another important component to change is the output location where the standard output and standard error files will be saved. If every script uses the same name for these output files, they will simply be overwritten whenever a new job is executed.

Depending on the number of jobs you are planning to submit, it can also be helpful to write a small program that actually generated the qsub bash scripts for you. An easy way to do this is to start with a base file that has almost all the information you need, excluding the names of the output text files and the different input arguments you want to supply. Then you can write a small script that simply fills these areas with the information you require. If you want some examples on how I do this, just let me know and I can send you some examples.

HPC will let you run roughly 100 different jobs simultaneously, though you are able to submit many many more jobs than that. A max of 10000 job submissions would be a reasonable limit, depending on the resources you require. If you do submit more than 100 jobs, excess jobs will simply join the cue and commence running after some of your other jobs have finished.

3.2 Multicore Processing - The Hard Way

Whilst the previous section described how to efficiently and easily parallelise your work, the bulk submission method is only suitable when individual tasks can be run independently. For many cases, this type of parallelisation is not possible. For these scenarios, parallelisation may still be feasible, as a few packages support multicore processing. In general, this is a much more involved and arduous task, and one that I don't believe R handles nicely when compared to other programming languages such as Python². Given the increased complexity associated with implementing many multiprocessing programs directly within R, it will not be covered in this guide, as it is assumed that if you have the programming proficiency to implement such a system, you should have little dramas migrating it to the HPC environment.

²Although it can also be a pain to implement multithreaded programs in Python!

4 Quick Guide

Command	Description/Example
cat	Display contents of a file cat <path_to_file>
cd	Change Directory cd ~ #change to home directory cd .. #move up one directory
cp	Copy file cp <path_original_file> <path_new_file>
ls	List files in current directy
man	Manual for a command man ls
module avail	List all of the available modules
module load	Load a specific module module load atg/R/3.4.1-foss-2016a
module purge	Remove all loaded modules
mv	Move a file mv <path_original_file> <path_new_file>
rm	Remove a file permanently rm <path_to_file>
qdel	Delete a specific job that was submitted to the queue qdel <job_id_number> #find number with qstat
qstat	View Running jobs qstat -u <your_QUT_username>
qsub	Submit a job to the queue # Interactive Job qsub -I -S /bin/bash -l walltime=HH:MM:SS,mem=XXXg,ncpus=YYY -N ZZZ # batch job qsub <path_to_batch_script>

5 Other Tips

This section will be updated every now and then with anything new I find. If you find an interesting tip, trick or some package specific information that you think others might benefit from, let me know and we can add it.

References

- [1] S. Tatham. (2018). Putty, [Online]. Available: <https://www.putty.org/>.
- [2] W. Wolber. (1992). Breast cancer wisconsin data set, UCI, [Online]. Available: [https://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+\(Original\)](https://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+(Original)).